

Report

Assignment 2 - MySQL

Group: 39

Students: Folke Jernbert

Introduction

This assignment was about, first, inserting a dataset into a MySQL database and then querying it for interesting patterns using MySQL-connector. The dataset was a GPS trajectory dataset recording time stamps of people doing activities in and around Beijing from April 2007 to August 2012. I developed the code for the assignment locally and then ssh-ed into a virtual machine provided by IDI for the implementation.

The *set-up* required first that I configured login credentials for the database. Then I used the provided template ``DbConnector.py`` for connecting to the database. I used this with little changes other than adding environment variables.

The logic for *inserting* the data was done in ``insert_data.py``. Three tables, 'User', 'Activity', and 'TrackPoint' were created. In inserting the tables, some data cleaning had to be done such as replacing altitudes set to -777 with NULL-values and ignoring non-valid transportation modes.

Once the database was filled with the dataset, I could start *querying*. Tasks 1 to 6 were performed simply by querying using SQL in the terminal. These queries were relatively simple, generally taking averages and counting occurrences. For tasks 7 to 10 Python was used with the connector. The implementation of these queries are provided as attachments with this hand-in. The final, most complicated task, 11, was solved using a complex SQL-query.

Overall I found the assignment quite enjoyable. A side effect of the project was to learn a couple of tricks in the terminal, like ``history``, ``!...`` and generally how powerful a VM can be when used in combination with git - this enabled me to develop locally and quickly push/pull when I wanted to run scripts. This was a very nice experience. The tasks themselves also gave a feeling of how to approach a dataset during exploratory analysis (EDA), finding counts and averages, checking for invalid points, interesting patterns related to features (for example altitude), and how to formulate interesting inquiries. I have done quite some EDA in Jupyter Notebooks before, using Pandas, but taking such a huge dataset, split up into a folder structure and with labels in a separate text-file, and inserting it into a database before querying it, seems a more, for lack of better words, "production-level solution". It seems the kind of thing data engineers are tasked with before handing over the cleaned database to analysts and data scientists. Getting an experience with such a task, was fun and good learning.

AI-technology like ChatGPT and CoPilot have become integrated and ingrained in my workflow, from project approach, code generation to evaluation. Besides this sparring partner, I worked alone with this project.

Results

Part 1

Showing tables in the database and top 10 entries for each of the tables.

```
mysql> SHOW TABLES;
+-----+
| Tables_in_activity_db |
+-----+
| Activity               |
| TrackPoint            |
| User                   |
+-----+
3 rows in set (0.00 sec)
```

```
mysql> SELECT * FROM Activity LIMIT 10;
+----+-----+-----+-----+-----+
| id | user_id | transportation_mode | start_date_time | end_date_time |
+----+-----+-----+-----+-----+
| 1  | 0       | NULL               | 2008-10-23 02:53:04 | 2008-10-23 11:11:12 |
| 2  | 0       | NULL               | 2008-10-24 02:09:59 | 2008-10-24 02:47:06 |
| 3  | 0       | NULL               | 2008-10-26 13:44:07 | 2008-10-26 15:04:07 |
| 4  | 0       | NULL               | 2008-10-27 11:54:49 | 2008-10-27 12:05:54 |
| 5  | 0       | NULL               | 2008-10-28 00:38:26 | 2008-10-28 05:03:42 |
| 6  | 0       | NULL               | 2008-10-29 09:21:38 | 2008-10-29 09:30:28 |
| 7  | 0       | NULL               | 2008-10-29 09:30:38 | 2008-10-29 09:46:43 |
| 8  | 0       | NULL               | 2008-11-03 10:13:36 | 2008-11-03 10:16:01 |
| 9  | 0       | NULL               | 2008-11-03 23:21:53 | 2008-11-04 03:31:08 |
| 10 | 0       | NULL               | 2008-11-10 01:36:37 | 2008-11-10 03:46:12 |
```

Checking that there in fact are activities with transportation mode not null:

```
mysql> SELECT * FROM Activity WHERE transportation_mode IS NOT NULL LIMIT 10;
+----+-----+-----+-----+-----+
| id | user_id | transportation_mode | start_date_time | end_date_time |
+----+-----+-----+-----+-----+
| 1188 | 10      | bus                | 2008-05-18 07:24:54 | 2008-05-18 08:03:44 |
| 1226 | 10      | taxi               | 2008-10-15 00:51:43 | 2008-10-15 01:10:59 |
| 1234 | 10      | taxi               | 2008-10-21 00:46:22 | 2008-10-21 00:58:38 |
| 1239 | 10      | taxi               | 2008-12-01 00:34:31 | 2008-12-01 00:53:34 |
| 2356 | 20      | walk               | 2011-08-27 06:13:01 | 2011-08-27 08:01:37 |
| 2368 | 20      | walk               | 2011-09-01 14:49:54 | 2011-09-01 15:05:11 |
| 2371 | 20      | walk               | 2011-09-03 04:42:02 | 2011-09-03 06:16:53 |
| 2372 | 20      | walk               | 2011-09-03 14:20:01 | 2011-09-03 15:00:54 |
| 2374 | 20      | walk               | 2011-09-04 14:32:55 | 2011-09-04 14:58:55 |
| 2376 | 20      | walk               | 2011-09-05 04:05:59 | 2011-09-05 04:36:16 |
```

```
mysql> SELECT * FROM TrackPoint LIMIT 10;
```

id	activity_id	lat	lon	altitude	date_days	date_time
1	1	39.984702	116.318417	492	39744.1201851852	2008-10-23 02:53:04
2	1	39.984683	116.31845	492	39744.1202546296	2008-10-23 02:53:10
3	1	39.984686	116.318417	492	39744.1203125	2008-10-23 02:53:15
4	1	39.984688	116.318385	492	39744.1203703704	2008-10-23 02:53:20
5	1	39.984655	116.318263	492	39744.1204282407	2008-10-23 02:53:25
6	1	39.984611	116.318026	493	39744.1204861111	2008-10-23 02:53:30
7	1	39.984608	116.317761	493	39744.1205439815	2008-10-23 02:53:35
8	1	39.984563	116.317517	496	39744.1206018519	2008-10-23 02:53:40
9	1	39.984539	116.317294	500	39744.1206597222	2008-10-23 02:53:45
10	1	39.984606	116.317065	505	39744.1207175926	2008-10-23 02:53:50

```
10 rows in set (0.00 sec)
```

```
mysql> SELECT * FROM User LIMIT 10;
```

id	has_labels
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	0
8	0
9	0

```
10 rows in set (0.00 sec)
```

Checking that some ids in fact has labels:

```
mysql> SELECT * FROM User WHERE has_labels = 1 LIMIT 10;
```

id	has_labels
10	1
20	1
21	1
52	1
53	1
56	1
58	1
59	1
60	1
62	1

```
10 rows in set (0.00 sec)
```

Part 2

Task 1

The number of users, activities and trackpoints there are in the dataset.

```
mysql> SELECT COUNT(*) AS UserCount FROM User;
+-----+
| UserCount |
+-----+
|         182 |
+-----+
1 row in set (0.03 sec)
```

```
mysql> SELECT COUNT(*) AS ActivityCount FROM Activity;
+-----+
| ActivityCount |
+-----+
|         16048 |
+-----+
1 row in set (0.01 sec)
```

```
mysql> SELECT COUNT(*) AS TrackPointCount FROM TrackPoint;
+-----+
| TrackPointCount |
+-----+
|        9681756 |
+-----+
1 row in set (1.53 sec)
```

Task 2

The average number of activities per user.

```
mysql> SELECT AVG(activity_count) AS AvgActivitiesPerUser
-> FROM (
->     SELECT user_id, COUNT(*) AS activity_count
->     FROM Activity
->     GROUP BY user_id
-> ) AS user_activity_counts;
```

```
+-----+
| AvgActivitiesPerUser |
+-----+
|          92.7630 |
+-----+
1 row in set (0.03 sec)
```

Task 3

Top 20 users with the highest number of activities.

```
mysql> SELECT user_id, COUNT(*) AS activity_count
-> FROM Activity
-> GROUP BY user_id
-> ORDER BY activity_count DESC
-> LIMIT 20;
```

```
+-----+-----+
| user_id | activity_count |
+-----+-----+
| 128     | 2102          |
| 153     | 1793          |
| 25      | 715           |
| 163     | 704           |
| 62      | 691           |
| 144     | 563           |
| 41      | 399           |
| 85      | 364           |
| 4       | 346           |
| 140     | 345           |
| 167     | 320           |
| 68      | 280           |
| 17      | 265           |
| 3       | 261           |
| 14      | 236           |
| 126     | 215           |
| 30      | 210           |
| 112     | 208           |
| 11      | 201           |
| 39      | 198           |
+-----+-----+
20 rows in set (0.00 sec)
```

Task 4

All users who have taken a taxi.

```
mysql> SELECT DISTINCT user_id
-> FROM Activity
-> WHERE transportation_mode = 'taxi';
```

user_id
10
58
62
78
80
85
98
111
128
163

10 rows in set (0.02 sec)

Task 5

The activity count of transportation modes.

```
mysql> SELECT transportation_mode, COUNT(*) AS activity_count
-> FROM Activity
-> WHERE transportation_mode IS NOT NULL
-> GROUP BY transportation_mode;
```

<u>transportation_mode</u>	activity_count
bus	199
taxi	37
walk	481
bike	262
car	419
run	1
train	2
subway	133
airplane	3
boat	1

Task 6

a)

The year with the most activities.

```
mysql> SELECT YEAR(start_date_time) AS year, COUNT(*) AS activity_count
-> FROM Activity
-> GROUP BY year
-> ORDER BY activity_count DESC
-> LIMIT 1;
```

year	activity_count
2008	5895

b)

The year with the most activity recorded hours.

```
mysql> SELECT YEAR(start_date_time) AS year,
-> SUM(TIMESTAMPDIFF(SECOND, start_date_time, end_date_time)) / 3600 AS total_hours
-> FROM Activity
-> GROUP BY year
-> ORDER BY total_hours DESC
-> LIMIT 1;
```

year	total_hours
2009	11612.4239

1 row in set (0.02 sec)

Although 2008 is the year with the most activity counts, 2009 is the year with the highest recorded number of hours.

Task 7

Total distance in km walked in 2008 by user 112.

See provided code for implementation (`task7.py`)

```
You are connected to the database: ('activity_db',)
-----

Total distance walked by user 112 in 2008: 115.47 km
-----
```

Task 8

The top 20 users who have gained the most altitude.

See `task8.py` for implementation details.

User ID	Total Gain (m)
128	2,258,775.31
153	2,215,779.55
4	1,174,001.00
41	920,172.20
62	846,984.90
3	818,103.00
163	807,991.02
85	785,333.50
144	737,268.68
30	602,820.00
39	528,484.00
25	491,420.00
84	477,892.00
167	427,787.01
0	427,120.00
2	413,386.00
140	369,913.61
37	366,855.50
126	345,512.82
34	325,712.90

Task 9

The number of invalid activities per user (listing only users who actually have invalid activities).

Code: `task9.py`

User ID	Invalid Activity Count		
0	101	33	2
1	45	34	88
2	98	35	23
3	179	36	34
4	219	37	100
5	45	38	58
6	17	39	147
7	30	40	17
8	16	41	201
9	31	42	55
10	50	43	21
11	32	44	32
12	43	45	7
13	29	46	13
14	118	47	6
15	46	48	1
16	20	50	8
17	129	51	36
18	27	52	44
19	31	53	7
20	20	54	2
21	7	55	15
22	55	56	7
23	11	57	16
24	27	58	13
25	263	59	5
26	18	60	1
27	2	61	12
28	36	62	249
29	25	63	8
30	112	64	7
31	3	65	26
32	12	66	6
		67	33
		68	139

69	6	104	97
70	5	105	9
71	29	106	3
72	2	107	1
73	18	108	5
74	19	109	3
75	6	110	17
76	8	111	26
77	3	112	67
78	19	113	1
79	2	114	3
80	6	115	58
81	16	117	3
82	27	118	3
83	15	119	22
84	99	121	4
85	184	122	6
86	5	123	3
87	3	124	4
88	11	125	25
89	40	126	105
90	3	127	4
91	63	128	720
92	101	129	6
93	4	130	8
94	16	131	10
95	4	132	3
96	35	133	4
97	14	134	31
98	5	135	5
99	11	136	6
100	3	138	10
101	46	139	12
102	13	140	86
103	24	141	1

141	1
142	52
144	157
145	5
146	7
147	30
150	16
151	1
152	2
153	557
154	14
155	30
157	9
158	9
159	5
161	7
162	9
163	233
164	6
165	2
166	2
167	134
168	19
169	9
170	2
171	3
172	9
173	5
174	54
175	4
176	8
179	28
180	2
181	14

Task 10

Users who have been in the Forbidden City of Beijing (within a radius of 100 meters of the coordinates lat: 39.916, lon: 116.397).

User ID
18
19
4
131

Task 11

The most used transportation mode of users who have registered it.

```
mysql> SELECT user_modes.user_id, user_modes.transportation_mode
-> FROM (
->     SELECT user_id, transportation_mode, COUNT(*) AS mode_count
->     FROM Activity
->     WHERE transportation_mode IS NOT NULL
->     GROUP BY user_id, transportation_mode
-> ) AS user_modes
-> JOIN (
->     SELECT user_id, MAX(mode_count) AS max_count
->     FROM (
->         SELECT user_id, transportation_mode, COUNT(*) AS mode_count
->         FROM Activity
->         WHERE transportation_mode IS NOT NULL
->         GROUP BY user_id, transportation_mode
->     ) AS counts
->     GROUP BY user_id
-> ) AS max_counts
-> ON user_modes.user_id = max_counts.user_id
->     AND user_modes.mode_count = max_counts.max_count
-> ORDER BY user_modes.user_id;
```

user_id	transportation_mode
10	taxi
20	bike
21	walk
52	bus
56	bike
58	taxi
58	car
58	walk
60	walk
62	bus
62	walk
64	bike
65	bike
67	walk
69	bike
73	walk
75	walk
76	car
78	walk
80	taxi
80	bike
81	bike
82	walk
84	walk
85	walk
86	car
87	walk
89	car
91	bus
91	walk
92	bus
92	walk
97	bike
98	taxi
101	car
102	bike
107	walk
108	walk
111	taxi
112	walk
115	car
117	walk
125	bike
126	bike
128	car
136	walk
138	bike
139	bike
144	walk
153	walk
161	walk
163	bike
167	bike
175	bus

54 rows in set (0.02 sec)

Discussion

In general I find the assignment went along quite smoothly. There were some initial hick-ups with opening the ports in the configuration file of MySQL, but this issue was quickly resolved. Syncing with GitHub was also in the beginning something of a pain point. My quota of Git LFS was used up, and so I couldn't transfer the data files using that solution. Resetting the commit took some time, but eventually I was able to add the datafiles to the .gitignore-file, and I used a file server to transfer the data files to the VM.

As for the tasks, I found it nice to brush up on some SQL querying from this spring. While the basics, like selecting, grouping by, and aggregating, feel quite comfortable, things was more difficult with constructing nested queries such as in task 11.

The concept of the MySQL-connector in Python was not entirely new to me, as I in the database course previous semester used an ORM with Flask in the main project. A difference, however, is that I in this task wrote SQL-queries directly into the Python script.

As briefly alluded to in the introduction, I found the tasks quite relevant for real-world scenarios, where typical tasks of the data handler are data cleaning, inserting data into a database, checking for basic statistics, and finding some basic insights like who has been in the Forbidden City. This is a nice skill to have, and I'm glad I've gotten some practice in that through this project.